

TRANSACTION MANAGEMENT

Part I

Sanit Sirisawatvatana and Sunisa Sathapornvajana

Chapter Objectives

- Function and importance of transactions.
- Properties of transactions.
- Concurrency Control
 - Locking
 - Deadlock
 - Timestamping
- Database Recovery
 - Transaction log
 - Checkpointing

Review

- Types of SQL statements
 - **DML**: Data Manipulation Language
 - SELECT, INSERT, UPDATE, DELETE statements, etc.
 - **DDL**: Data Definition Language
 - CREATE, ALTER, DROP, TRUNCATE statements, etc.
 - **DCL**: Data Control Language
 - GRANT and REVOKE statements
 - **TCL**: **Transaction Control Language**
 - COMMIT, ROLLBACK and SAVEPOINT statements

Transaction

- **Definition:**
 - Action, or series of actions, carried out by user or application, which reads or updates contents of database.
- Transaction
 - Is a logical unit of work on the database
 - May be an entire (or a part of) program or a single statement that involve any number of operations on the database
 - Transforms database from one consistent state to another
- Database operations
 - Read/Write

Example: Transaction

```
read(staffNo = x, salary)
salary = salary * 1.1
write(staffNo = x, new_salary)
```

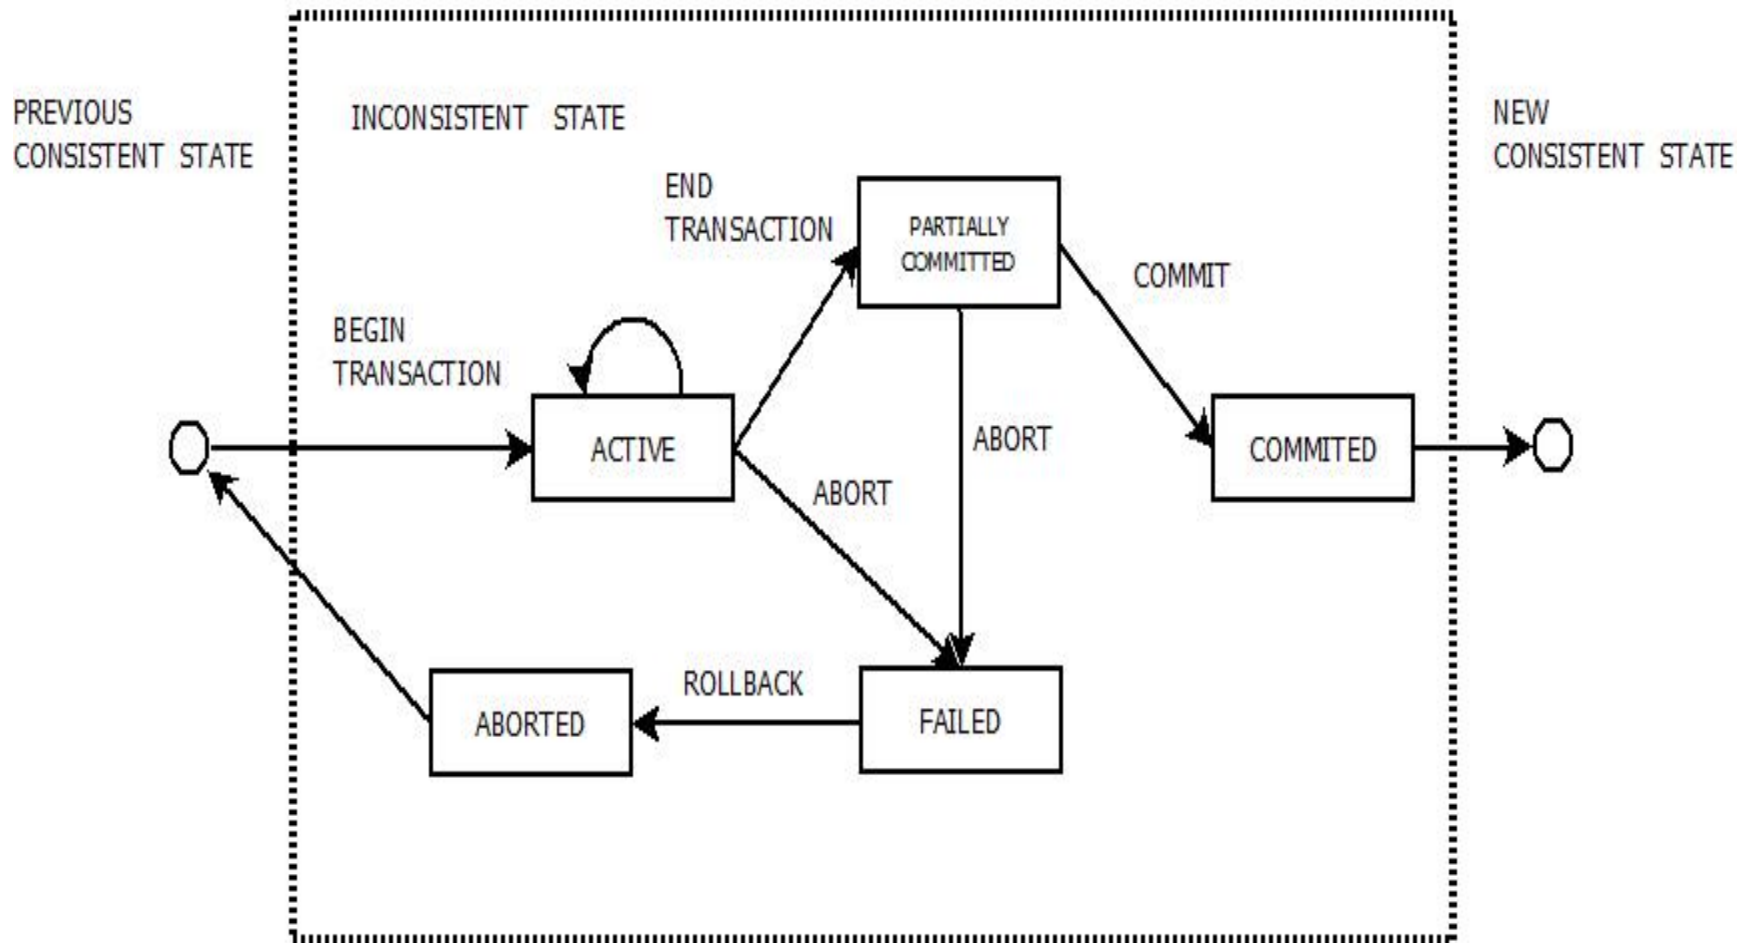
(a)

```
delete(staffNo = x)
for all PropertyForRent records, pno
begin
  read(propertyNo = pno, staffNo)
  if (staffNo = x) then
    begin
      staffNo = newStaffNo
      write(propertyNo = pno, staffNo)
    end
  end
end
```

(b)

State Transition Diagram for a Transaction

Transaction



States of a Transaction

- **ACTIVE**
 - Executes database and non-database operations
- **COMMITTED**
- **ABORTED**
- **PARTIALL COMMITED**
 - Occurs after the final statement has been executed
 - It may be found that transaction has violated an integrity constraint
- **FAILED**
 - Occurs when transaction cannot be committed or transaction is aborted
 - The user might abort the transaction
 - The concurrency control protocol might abort the transaction to ensure serializability

Outcomes of a Transaction

Two outcomes:

- **COMMITTED**
 - Transaction **completes** successfully
 - The database reaches a **new consistent state**
- **ABORTED**
 - Transaction does **not execute successfully**
 - The database **must be restored** to the consistent state (**previous consistent state**) before it started (**Rolled back** or **undone**)
- A **committed transaction** cannot be aborted.
- An **aborted transaction** that is rolled back can be restarted later and may execute successfully and commit in later time.

Properties of Transactions (ACID)

All transactions should have basic four properties:

- **Atomicity**
 - An **Indivisible unit**
 - All or Nothing
- **Consistency**
 - Database transforms from one consistent state to another consistent state (consistent **before** and **after** transaction)
- **Isolation**
 - Executes independently of one another (**not interfere** with each other).
- **Durability**
 - The completed transaction is **permanently recorded** in database.
 - The completed transaction must **not be lost from failure**.

Example: Transaction

A banking application is the classic example of why transactions are necessary. Imagine a bank's database with two tables: checking and savings. To move \$200 from Jane's checking account to her savings account, you need to perform at least three steps:

1. Make sure her checking account balance is greater than \$200.
2. Subtract \$200 from her checking account balance.
3. Add \$200 to her savings account balance.

The entire operation should be wrapped in a transaction so that if any one of the steps fails, any completed steps can be rolled back. You start a transaction with the `START TRANSACTION` statement and then either make its changes permanent with `COMMIT` or discard the changes with `ROLLBACK`. So, the SQL for our sample transaction might look like this:

```
1 START TRANSACTION;  
2 SELECT balance FROM checking WHERE customer_id = 10233276;  
3 UPDATE checking SET balance = balance - 200.00 WHERE customer_id = 10233276;  
4 UPDATE savings SET balance = balance + 200.00 WHERE customer_id = 10233276;  
5 COMMIT;
```

ACID Transaction

**A**

Atomicity requires that each transaction be "all or nothing": if one part of the transaction fails, then the entire transaction fails, and the database state is left unchanged.

**C**

The consistency property ensures that any transaction will bring the database from one valid state to another.

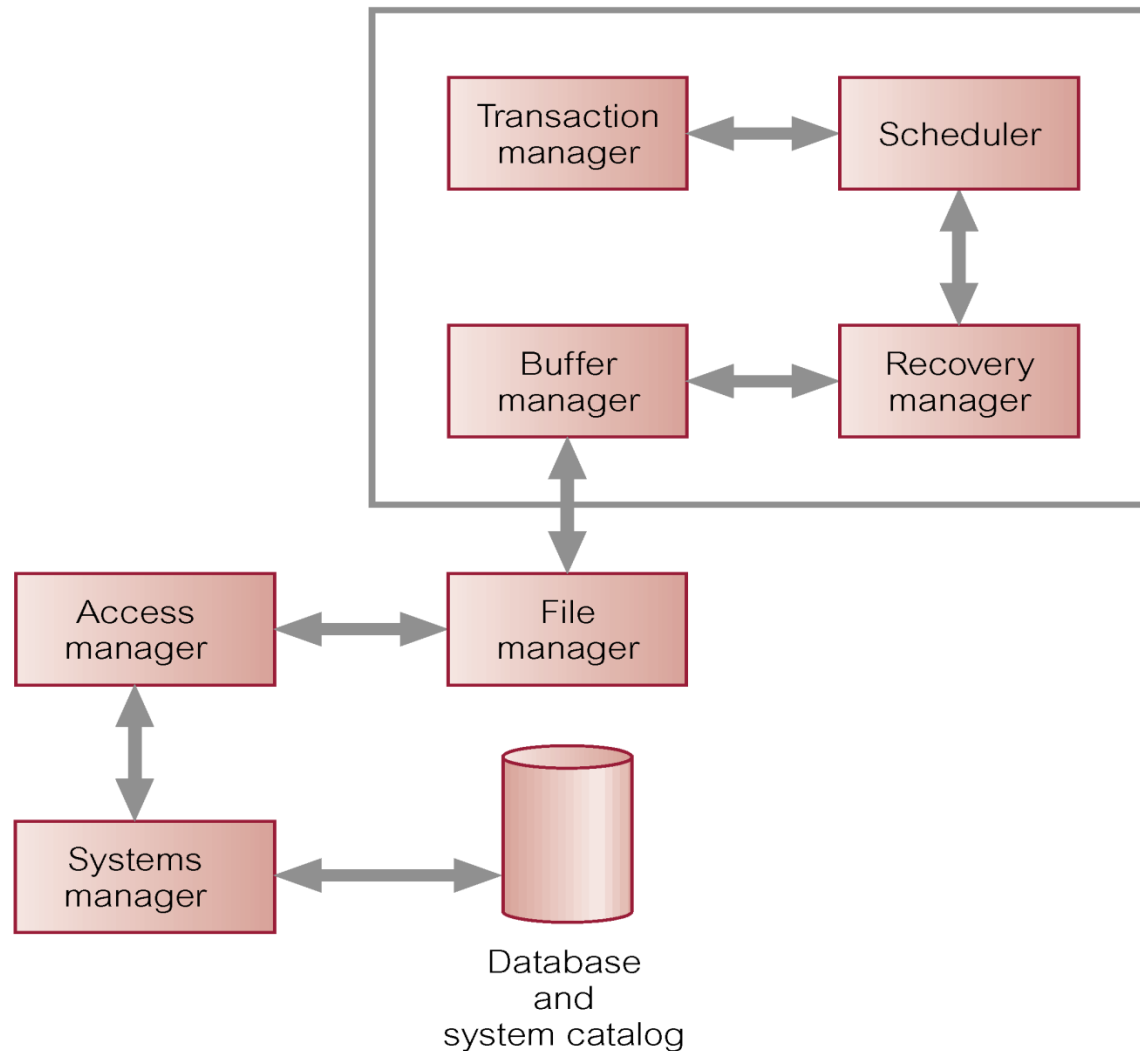
**I**

The isolation property ensures that the concurrent execution of transactions results in a system state that would be obtained if transactions were executed serially, i.e., one after the other.

**D**

The durability property ensures that once a transaction has been committed, it will remain so, even in the event of power loss, crashes, or errors.

DBMS Transaction Subsystem



DBMS Transaction Subsystem

- **Concurrency Control**
 - To maximize concurrency without allowing concurrently executing transactions to interfere with each other
- **Database Recovery**
 - To ensure that database is consistent when a failure occurs during the transaction

Concurrency Control

- **Definition:**
 - Process of managing simultaneous operations on the database without having them interfere with one another.
- when two or more users are **accessing database simultaneously**
 - **All users are only reading**
 - => No way to interfere one another
 - **At least one is updating data**
 - => May be interfere one another
 - => Can result in inconsistencies

Problems in Concurrency

Potential problems caused by concurrency:

- Lost update problem
- Uncommitted dependency (**dirty read**)
- Inconsistent analysis problem

The Lost Update Problem

- Successfully completed update is overridden by another user.

Time	T_1	T_2	bal_x
t_1		begin_transaction	100
t_2	begin_transaction	read(bal_x)	100
t_3	read(bal_x)	$bal_x = bal_x + 100$	100
t_4	$bal_x = bal_x - 10$	write(bal_x)	200
t_5	write(bal_x)	commit	90
t_6	commit		90

- At t_4 T_2 Loss the updated bal_x 200
- This can avoid by preventing T_1 from reading bal_x until after update.

Uncommitted Dependency Problem

- Occurs when one transaction can see intermediate results of another transaction before it has committed.

Time	T ₃	T ₄	bal _x
t ₁		begin_transaction	100
t ₂		read(bal _x)	100
t ₃		bal _x = bal _x + 100	100
t ₄	begin_transaction	write(bal _x)	200
t ₅	read(bal _x)	⋮	200
t ₆	bal _x = bal _x - 10	rollback	100
t ₇	write(bal _x)		190
t ₈	commit		190

- T4 updates bal_x to £200 but it aborts
- T3 has read new value of bal_x (£200) giving a new balance of £190, instead of £90.
- This problem can be avoided by preventing T3 from reading bal_x until after T4 commits or aborts.

Inconsistent Analysis Problem

- Occurs when transaction reads several values but second transaction updates some of them during execution of first.

Time	T ₅	T ₆	bal _x	bal _y	bal _z	sum
t ₁		begin_transaction	100	50	25	
t ₂	begin_transaction	sum = 0	100	50	25	0
t ₃	read(bal _x)	read(bal _x)	100	50	25	0
t ₄	bal _x = bal _x - 10	sum = sum + bal _x	100	50	25	100
t ₅	write(bal _x)	read(bal _y)	90	50	25	100
t ₆	read(bal _z)	sum = sum + bal _y	90	50	25	150
t ₇	bal _z = bal _z + 10		90	50	25	150
t ₈	write(bal _z)		90	50	35	150
t ₉	commit	read(bal _z)	90	50	35	150
t ₁₀		sum = sum + bal _z	90	50	35	185
t ₁₁		commit	90	50	35	185

- The problem can be avoided by preventing T₆ from reading bal_x and bal_z until after T₅ completed updates.

Inconsistent Retrieval Problems

- Interference causes inconsistency among multiple retrievals of a subset of data
 - Incorrect summary (**Inconsistent analysis problem**)
 - Non-repeatable (or **fuzzy**) read
 - Phantom read

Serializability and Recoverability

- When multiple transactions run concurrently, there is a possibility that the database may be left in an **inconsistent state**.
- **Serializability** is a concept that helps to identify which non-serial schedules are correct and will maintain the consistency of the database.
- **Schedule**
 - A sequence of the operations by a set of concurrent transactions that prevents the order of the operations in each of the individual transactions.

Serial Schedule vs. Non-serial Schedule

- **Serial schedule**
 - Is always a serializable schedule.
 - A schedule where operations of each transaction are executed consecutive without any interleaved operations from other transactions.
 - A transaction only starts when the other transaction finished executed.
- **Non-serial schedule**
 - A schedule where the operations from a set of current transactions **are interleaved**.
 - Is said to be serializable schedule, if it is equivalent to the serial schedule of those n transactions.

Serializability

- **The objective of serializability**
 - Is to find non-serial schedules that allow transactions to execute concurrently without interfering with one another, and thereby produce a database state that could be produced by a serial execution.
- **In serializability, the ordering of read/writes is important:**
 - If two transactions **only read a data item**, they do **not conflict** and **order is not important**.
 - If two transactions **either read or write** completely **separate data items**, they **do not conflict** and **order is not important**.
 - If one transaction **writes a data item** and **another reads or writes same data item**, **order of execution is important**.

Recoverability

- Serializability identifies schedules that maintain the consistency of database (none of the transactions in the schedule fails)
- Recoverability of transactions within a schedule:
 - If a transaction fails, the **atomicity property** requires that we undo the effects of transactions.
 - The **durability property** states that once a transaction commits, its changes cannot be undone (without running another, compensating, transaction).

Unrecoverable Schedule

- G is **unrecoverable schedule**, because T2 read the value of A written by T1, and committed.
- T1 later aborted, therefore the value read by T2 is wrong, but since T2 committed.

$$G = \begin{bmatrix} T1 & T2 \\ R(A) & \\ W(A) & \\ & R(A) \\ & W(A) \\ & Com. \\ Abort & \end{bmatrix}$$

Recoverable Schedule

- A schedule where, for each pair of transactions T_i and T_j , if T_j reads a data item previously written by T_i , then the commit operation of T_i precedes the commit operation of T_j .

$$F = \begin{bmatrix} T1 & T2 \\ R(A) & \\ W(A) & \\ & R(A) \\ & W(A) \\ Com. & \\ & Com. \end{bmatrix} \quad F2 = \begin{bmatrix} T1 & T2 \\ R(A) & \\ W(A) & \\ & R(A) \\ & W(A) \\ Abort & \\ & Abort \end{bmatrix}$$

- F is **recoverable** because $T1$ commits before $T2$, that makes the value read by $T2$ corrects, Then $T2$ can commit itself.
- In $F2$, if $T1$ aborted, $T2$ has to abort because the value of A it read is incorrect

Concurrency Control Techniques

- Two concurrency control techniques
 - Locking
 - Timestamping
- Both are **conservative** (or **pessimistic**) approaches
 - They cause transaction to be delayed when they conflict with other transactions
- **Optimistic** approaches based on:
 - Transaction conflict **is rare**. So they allow transactions to run unsynchronized and check for conflicts only when the transaction commits at the end

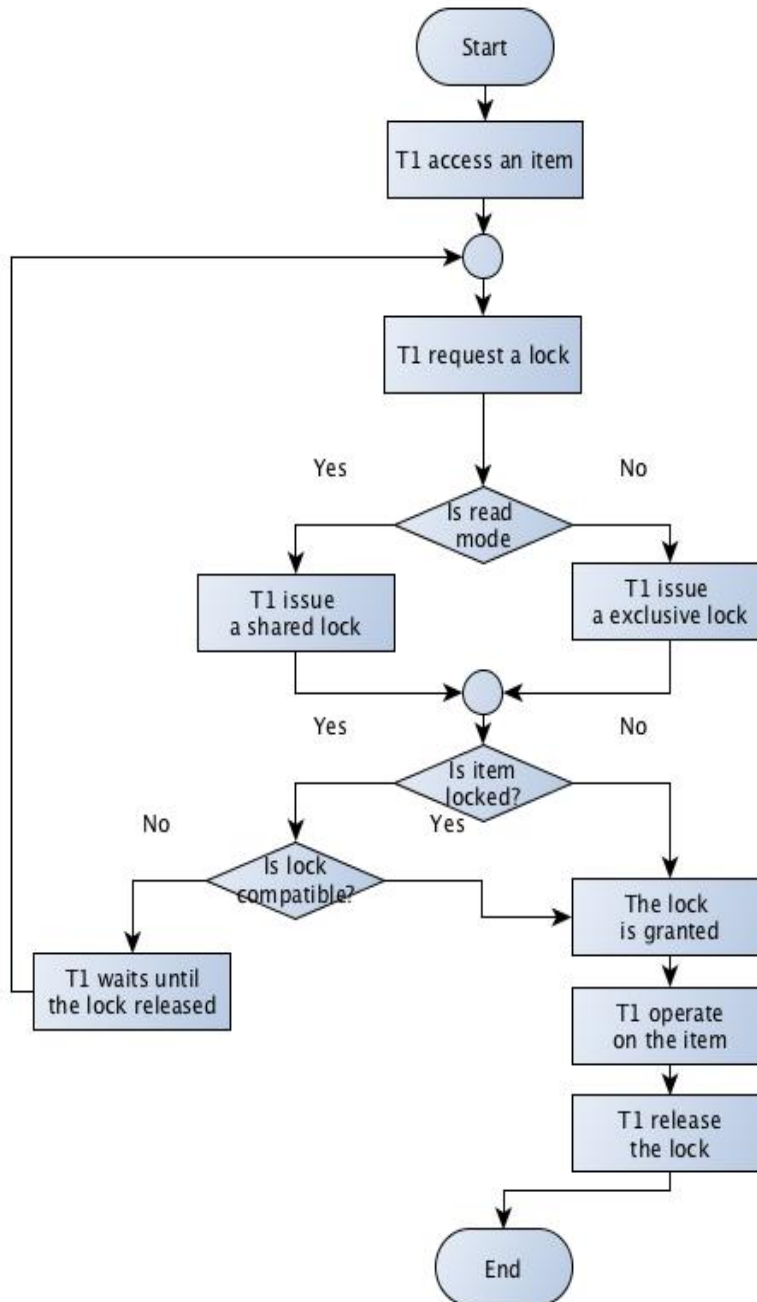
Locking Methods

- **Definition:**
 - A **procedure** used to control **concurrency access** to data. Transaction uses locks to **deny access** to other transactions to **prevent incorrect results**.
- **Locking methods** are the most widely used approach to ensure **serializability of concurrent transactions**.
- Transaction must request **a lock** on a data item before it read or write the data item.
- The lock **prevents** another transaction from modifying or reading the item.
- The other transaction must wait until the lock is released.

Types of a lock

- Two types:
 - **Shared lock** (read lock)
 - Is used for **read-only** mode (**reading purpose**)
 - Transaction requests shared lock on a data item in order **to read its content**.
 - Updating the data is not allowed
 - **Exclusive lock** (write lock)
 - Is used for **write mode**
 - Only the transaction that requests a lock can read and update the data item
 - The other transactions can not read , write or lock the data and must wait until the lock is released

Flowchart of locking



Lock Matrix

T1 issue a lock at t = 0	
T2 issue a lock at t = 10	
Shared	Exclusive
Shared	Both are granted shared locks Both only Read
Exclusive	Only T1 is granted a shared lock T2 wait for a lock T1 only Read

Some systems allow transaction to **upgrade** shared (read) lock to an exclusive lock, or **downgrade** exclusive lock to a shared lock.

Locking Granularity

